

TU DORTMUND

INTRODUCTORY CASE STUDIES

Project 3: Regression analysis for price of VW cars

Lecturers:

Prof. Dr. Sonja Kuhnt

Dr. Paul Wiemann

Dr. Birte Hellwig

M. Sc. Hendrik Dohme

Author: Tadeo Hepperle

Group number: 2

Group members: Prem Kant Shekhar, Minjae Ok, Ishan Singh
Dhapola, Fyalisia Amanda Putri, Tadeo Hepperle

January 17, 2022

Contents

1	Introduction	3
2	Problem statement	3
2.1	Description of the dataset	3
2.2	Objectives of the report	4
3	Statistical Methods	4
3.1	classical linear regression	5
3.2	dummy coding of categorical variables	8
3.3	model diagnostics	8
3.4	best subset selection	8
3.4.1	adjusted R^2	9
3.4.2	Akaike information criterion (AIC)	9
3.4.3	Bayesian information criterion (BIC)	10
3.5	inference and linear models	10
4	Results	11
4.1	Descriptive Statistics	11
4.2	Determining the dependent variable	12
4.3	best subset selection	13
5	Summary	14
	Bibliography	15
	Appendix	16
A	Additional figures	16

1 Introduction

2 Problem statement

We analyzed a dataset containing prices and other features of 438 cars from the german car brand VW to determine which factors influence the price of a car and to what extent they do so. This makes it possible to predict the value of a car by knowledge about its features.

2.1 Description of the dataset

The dataset used in this report is a slice of a bigger dataset from kaggle.com and contains data about 438 VW cars, which were originally scraped in 2020 from the british car internet platform "Exchange and Mart" (Exc) (<https://www.exchangeandmart.co.uk/>). The slicing was done in a way to only include the VW models "Passat", "T-Roc" and "Up". For each of the 438 cars information on the following variables is provided:

- the **price** a car is selling for in 1000 GBP (£), metric variable
- the **year** a car was first registered, metric variable
- the **mileage** as the total distance a car has been driven in 1000 miles, metric variable
- the **fuel consumption** in mpg, the number of miles a car can drive with one gallon of fuel, metric variable
- the **fuel type** a car uses, categorical variable with 3 levels: "Diesel", "Hybrid", "Petrol"
- the **engine size** as the volume of fuel in liters and air a car can fit in its engines cylinders, metric variable
- the annual **tax**, also known as "Vehicle Excise Duty" that has to be paid for the car annually
- the type of **transmission** a car has, categorical variable with 3 levels: "Manual", "Semi-Auto", "Automatic"
-
- the **model** of the car, categorical variable with 3 levels: "Passat", "T-Roc", "Up"

From those 9 original variables, three more are computed:

- the **logprice** as the natural logarithm of the **price** of a car
- the **fuel consumption** in liters per 100 kilometers, computed from the fuel consumption in mpg as $\frac{282.48}{mpg}$
- the **age** of the car in 2020 when the data was taken, computed as $2020 - year$

From here on, when fuel consumption is mentioned, the liters per 100 kilometers value is meant. The quality of the data seems good, "Exchange and Mart" is a large and well known site and no data is missing. Although we have to rely on the data given to us and cannot check if for example the mileage and tax are truly correct.

2.2 Objectives of the report

The goal of this report is to predict the price of a car as accurately as possible. To achieve this we fit a multitude of linear regression models and using best subset selection will choose as the final model the one with the best indicators, namely the BIC (Bayesian information criterion). Interactions between the 8 predictors or nonlinear relationships will not be considered.

Also we will figure out if it makes more sense to predict the price directly through a linear regression model or if it is better to use the logprice as the dependent variable. After making predictions this can be retransformed to normal prices so no information is lost, but it could be that the logarithmic price fits the data better. Sadly the range of car models in the data is quite limited and therefore we will only be able to predict prices of those models later on with our regression models.

3 Statistical Methods

We give a brief overview about the math behind multiple linear regression, how to select the best model in subset selection via certain indicators and how to assess how good a regression model fits the data.

3.1 classical linear regression

Multiple linear regression is a method for predicting a metric variable y on the basis of k metric variables x_1, \dots, x_k . A linear regression model consists of k coefficients that have to be fitted to the data and can be represented by the following formula:

$$y = \beta_1 x_1 + \dots + \beta_k x_k + \epsilon$$

y is also called dependent variable or outcome, while x_1 to x_k are called the independent variables or predictors. Since the coefficients β_1, \dots, β_k are just multiplied by the predictors x_1, \dots, x_k , a linear regression can only detect linear relationships in the data. In the formula ϵ represents the error term, because it is likely that even the best linear combination of the multiple $\beta_j x_j$ will not add up to the real value of y . But what do we mean by "the best linear combination"? Typically the coefficients in linear regression are estimated by minimizing the residual square sum (RSS). If we have all x_1 to x_k and β_1 to β_k we can estimate y as \hat{y} :

$$\hat{y} = \beta_1 x_1 + \dots + \beta_k x_k$$

The error term ϵ is missing in this equation. $y = \hat{y} + \epsilon$ or $\epsilon = y - \hat{y}$. Suppose we have a data set consisting of n objects i with values y_i and $x_{1,i}$ to $x_{k,i}$. Each object has its own error term ϵ_i , sometimes we overestimate y with \hat{y} , sometimes we underestimate it. After fitting the model coefficients by minimizing the RSS, the mean error ($\sum_{i=1}^n \epsilon_i$) will be zero.

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (\epsilon_i)^2$$

To estimate the coefficients, we introduce some matrix notation. Let Y be a $n \times 1$ column vector containing y_1 to y_n . X is a $n \times k$ -matrix consisting of n row vectors that contain the values $x_{1,i}$ to $x_{k,i}$ for each object. Then we can define \mathcal{B} as a $k \times 1$ vector containing the coefficients β_1, \dots, β_k and we define \mathcal{E} as a $n \times 1$ column vector containing ϵ_1 to ϵ_n . In this setting we can then write the formula as the following:

$$Y = X\mathcal{B} + \mathcal{E}$$

Here the RSS depends on \mathcal{B} can be rewritten as follows (Fahrmeir et al., 2013, p. 105):

$$\begin{aligned}
RSS(\mathcal{B}) &= \mathcal{E}^T \mathcal{E} \\
&= (Y - X\mathcal{B})^T (Y - X\mathcal{B}) \\
&= Y^T Y - B^T X^T Y - Y^T X\mathcal{B} + \mathcal{B}^T X^T X\mathcal{B} \\
&= Y^T Y - 2Y^T X\mathcal{B} + \mathcal{B}^T X^T X\mathcal{B}
\end{aligned} \tag{1}$$

Taking the derivative with respect to \mathcal{B} now yields:

$$\frac{\partial}{\partial \mathcal{B}} RSS(\mathcal{B}) = -2X^T Y + 2X^T X\mathcal{B}$$

It can be shown that the second derivative is positive and therefore we can find a minimum for $\mathcal{B} = \hat{\mathcal{B}}$ by setting the first derivative to zero (Fahrmeir et al., 2013, p. 106). $\hat{\mathcal{B}}$ represents the least squares estimator for \mathcal{B} .

$$\begin{aligned}
-2X^T Y + 2X^T X\hat{\mathcal{B}} &= 0 \\
2X^T X\hat{\mathcal{B}} &= 2X^T Y \\
\hat{\mathcal{B}} &= (X^T X)^{-1} X^T Y
\end{aligned} \tag{2}$$

So this is how we calculate our estimated coefficients $\hat{\beta}_1, \dots, \hat{\beta}_k$ as entries in the vector $\hat{\mathcal{B}}$. This classical linear model is missing an intercept though. Setting all predictors to zero would always yield $\hat{y} = 0$, no matter the coefficients. This is obviously not an optimal estimate. To accomodate for this, often a predictor x_0 is introduced that gets a value of 1 assigned for each object in the data. As a part of the predictor matrix X the method described above will then estimate a coefficient $\hat{\beta}_0$, known as the intercept.

A coefficient $\hat{\beta}_i$ can be interpreted in the following way: Holding all other predictor variables constant, an increase of one unit in x_i will on average result in an increase of \hat{y} by $\hat{\beta}_i$. The intercept β_0 represents the expected value of y when $x_1 = \dots = x_k = 0$. Sometimes logarithmic transformations are applied to the variables before fed into in the linear regression. When the outcome y is actually a proxy for $\ln(y_{original})$, than a one unit increase in x_i leads to a $\hat{\beta}_i$ increase in y which means a multiplication of $y_{original}$ by $10^{\hat{\beta}_i}$. For small $\hat{\beta}_i$ this means: a one unit increase in x_i leads to a $100 \cdot \hat{\beta}_i\%$ increase in $y_{original}$.

To assess how good a model is, a statistic R^2 can be calculated. It represents how much variance in y can be explained by x_1, \dots, x_k , or in other words how much variance y

shares with \hat{y} . Therefore it can be calculated as 1 minus the fraction between RSS and TSS where TSS is the empirical variance of y multiplied by n (James et al., 2013, p. 234).

$$\begin{aligned}
R^2 &= 1 - \frac{RSS}{TSS} \\
&= 1 - \frac{RSS}{\sum (y_i - \bar{y})^2} \\
&= 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2} \\
&= \frac{\sum (\hat{y}_i - \bar{y})^2}{\sum (y_i - \bar{y})^2}
\end{aligned} \tag{3}$$

It can also be calculated as the squared correlation between the real outcome y and the predicted outcome \hat{y} (Fahrmeir et al., 2013, p. 113).

$$R^2 = r_{y\hat{y}}^2$$

Therefore R^2 can take values between 0 and 1. A greater R^2 means a more accurate prediction. What a good R^2 value is, highly depends on the application though, while in physics often values close to 1 can be found, in "biology, psychology, marketing, and other domains" values below $R^2 \approx 0.10$ are to be expected according to James et al. (2013, p. 70).

A linear regression should only be used if some assumptions are fulfilled. First, the number of observations should be greater than the number of predictors. The residuals should be independent of the predicted values. For this, a scatter plot for \hat{y} vs. ϵ should not show a specific pattern. Their variance should not change as predicted values increase or decrease. This is also known as equal variance or homoscedasticity (Prabhakaran). Also the correlation between residuals and each predictor should be zero. This however is fulfilled when minimizing the RSS, because otherwise, the predictors "know" something about the residuals, which means they know more about y and are therefore not optimal. In addition to that the residuals should be approximately normally distributed which can be checked by looking at a QQ-plot between quantiles of standardized residuals and the standard-normal distribution (Prabhakaran).

3.2 dummy coding of categorical variables

The question arises how to use categorical variables as predictors in multiple linear regression. This can be resolved by a coding system. A coding system creates for every categorical variable a number of derived metric variables that can then be used in linear regression. One such coding system is dummy coding. citeXXXXX Assuming we have a categorical variable x_i with $m \geq 2$ levels with labels l_1, \dots, l_m . We can then create a dichotomous variable x_{l_j} for each label $l_j \in \{l_2, \dots, l_m\}$, such that:

$$x_{l_j} = \begin{cases} 1, & \text{if } x_i = l_j \\ 0, & \text{otherwise} \end{cases}$$

In this way all information about the categorical variable x_i is contained in the variables x_{l_2}, \dots, x_{l_m} . A datapoint with label $x_i = l_1$ can be recognized by having $x_{l_2} = \dots = x_{l_m} = 0$, l_1 is also called the baseline according to James et al. (2013, p. 86). Those newly created numerical variables can then be used in regression analysis as predictors.

3.3 model diagnostics

Assessing how well a linear model fits the data, is aided by some graphical and numerical methods. First, we can look at R^2 as a measure of goodness of fit. Also it is important to check the assumptions of a linear regression. For this some plots can help (Und)):

- Q-Q-plot of standardized residuals and a standard normal distribution check if residuals are approximately normally distributed
- Q-Q-plot of y and \hat{y} to check if they follow approximately the same distribution
- scale-location-plot: a scatter plot of \hat{y} on the x-axis and standardized residuals on the y-axis to see if they are evenly distributed along \hat{y} .

3.4 best subset selection

The more predictors go into a linear regression model, the more accurate it will perform on the training data and the greater its R^2 value will be. But taking more predictors into account is not always useful. It makes the model harder to interpret and does not highlight which variables are actually important for the outcome. Moreover a lot of

times variance in y that can be explained by an additional predictor is already explained by other predictors, such that the gain in variance explanation through R^2 is marginal and could even just be a product of overfitting. XXXXjameschapter2:Rsquared on train data is higher than on test data. To determine which predictors should actually be taken into account there is a method called "best subset selection" as described by James et al. (2013, p. 227).

To perform best subset selection on data with k predictors, for each $i \in 0, 1, \dots, k$ all possible linear models with a number of exactly i predictors are computed. There are $\binom{k}{i}$ different models for each i . Among those models for every i the model with least RSS is selected. This results in $k + 1$ models (including the empty model) that need to be considered to choose the best model. Finally from those $k + 1$ models the one with the best score on some indicator is chosen. possible indicators can be:

3.4.1 adjusted R^2

While R^2 explains how much variance in the training data can be explained by the predictors it is an overestimation for the performance on actual test data. Also it gets only larger with more predictors which might be a problem. Therefore an *Adjusted* R^2 can be calculated (James et al., 2013, p. 234).

$$AdjustedR^2 = 1 - \frac{RSS/(n - k - 1)}{TSS(n - 1)}$$

The term above the division line contains k and punishes models with more predictors. When used as an indicator in best subset selection the model with the greatest adjusted R^2 should be chosen.

3.4.2 Akaike information criterion (AIC)

According to James et al. (2013, p. 234) the AIC can be calculated using the following formula, where k denotes the number of predictors and $\hat{\sigma}^2$ is the variance of ϵ as $VAR(y - \hat{y})$ when looking at the model using all available predictors.

$$AIC = \frac{1}{n}(RSS + 2k\hat{\sigma}^2)$$

Please note that for simplicity normalizing constants have been left out as they do not matter when comparing two models (James et al., 2013, p. 234). A better model is characterized by a lower AIC.

3.4.3 Bayesian information criterion (BIC)

The Bayesian information criterion can be computed quite similarly to the AIC:

$$BIC = \frac{1}{n}(RSS + \ln(n)k\hat{\sigma}^2)$$

A low BIC characterizes a good model. The only difference is that the BIC uses the natural logarithm of n in the formula instead of 2. Because $\ln(n) > 2$ for $n \geq 8$ and we usually have more than 8 datapoints in our data, the BIC penalizes more predictors more heavily than the AIC and will therefore result in models with less predictors compared to the AIC when used in best subset selection (James et al., 2013, p. 234).

3.5 inference and linear models

To check if a linear regression model with k predictors can predict a significant portion of the variance in the dependent variable an F statistic can be calculated (James et al., 2013, p. 76). Under the assumption that the H_0 (=model cannot predict variance in outcome) is true, F follows an $F_{a,b}$ distribution with $a = n - k$ and $b = k - 1$.

$$F = \frac{(TSS - RSS)/k}{RSS/(n - k - 1)}$$

Each coefficient $\hat{\beta}_j, j \in 1, \dots, k$ can also be tested for statistical significance. For this, according to James et al. (2013, p. 67), we can compute a t statistic t_j by dividing the difference between $\hat{\beta}_j$ and the coefficient assumed under the null hypothesis (β_{H_0}) by the standard error of $\hat{\beta}_j$:

$$t_j = \frac{\hat{\beta}_j - \beta_{H_0}}{SE(\hat{\beta}_j)}$$

Under H_0 this would follow a t-distribution with $n-k-1$ degrees of freedom, which can then be compared to the $1-\frac{1}{2}\alpha$ -quantile of said distribution as a critical value. The standard error $SE(\hat{\beta}_j)$ can be computed with the following formula (James et al., 2013,

p. 66), where $\hat{\sigma}^2$ is the residual standard error that can be calculated from the residuals.

$$SE(\hat{\beta}_j) = \sqrt{\frac{\hat{\sigma}^2}{\sum_{i=1}^n (x_i - \bar{x})^2}}$$

$$\hat{\sigma}^2 = \sqrt{RSS/(n - 1 - k)}$$

This also allows for computing a $1-\alpha$ -confidence interval for each $\hat{\beta}_j$, that signify that $\hat{\beta}_j$ differs significantly from 0 if they do not contain the 0. a $1-\alpha$ -confidence interval for a parameter $\hat{\beta}_j$ has the property, that if we would take many random samples from the same population, $1-\alpha \cdot 100\%$ of confidence intervals will contain the true unknown parameter β_j (James et al., 2013, p. 66).

4 Results

First some descriptive statistics about the dataset is provided, then we evaluate which outcome will be used in the regression (price vs logprice) and finally the best model is computed and evaluated.

4.1 Descriptive Statistics

Table 1 provides a Five-number summary, mean and standard deviation for the relevant matrix variables used in the regression.

Table 1: Relevant Metric variables in the dataset

Rental Price	price	logprice	age	mileage	fuel consumption	tax	engineSize
mean	14.68	2.54	3.79	25.11	5.11	96.80	1.47
sd	7.75	0.55	1.96	25.04	1.19	61.65	0.42
minimum	3.50	1.25	1.00	1.20	1.70	0.00	1.00
Q1	7.78	2.05	2.00	6.05	4.40	20.00	1.00
Q2 (median)	12.00	2.48	4.00	17.53	4.70	145.00	1.50
Q3	20.99	3.04	5.00	33.37	5.60	145.00	2.00
maximum	38.99	3.66	15.00	138.57	8.69	265.00	2.00

In addition to that we have 3 categorical variables with 3 levels each: model, fuel consumption and transmission. Of the 438 cars in total, there were 161 Passat, 127 T-Roc and 150 Up. The majority of cars had manual transmission (320), 66 cars had

Semi-Auto and 52 automatic transmission. Most Cars used Petrol (256) or Diesel (169) with just 13 Hybrid cars that all where from the model Passat and had either automatic (4) or Semi-Auto transmission (9). The 150 Up models all had just manual transmission and ran on Petrol.

4.2 Determining the dependent variable

Two full linear regression models with all 8 predictor variables (age, mileage, fuel consumption, tax, engineSize, model, fuelType and transmission) were fitted to predict price and logprice. For the categorical variables dummy coding was utilized, where Passat was the reference category for model, Diesel for fuelType and Manual for transmission. This will stay constant for all following regressions in this report.

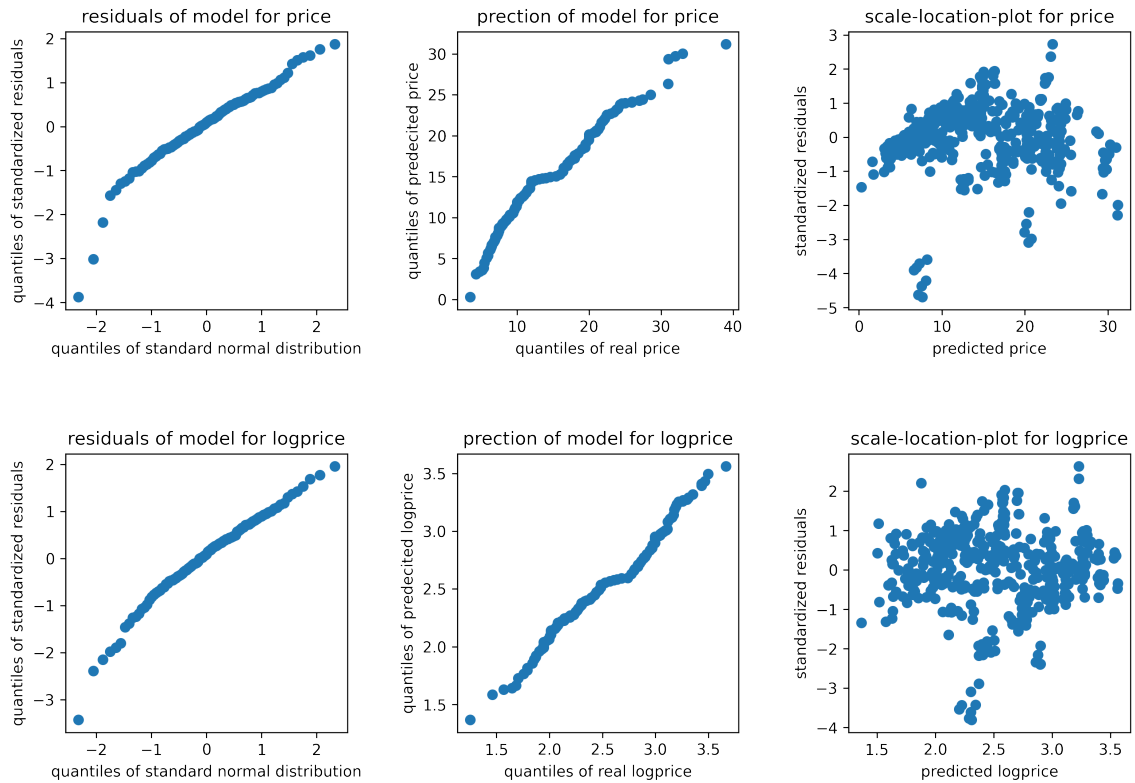


Figure 1: q-q-plots comparing quantiles of standardized residuals with standard normal distribution and predicted values with real values, scale-location-plots for relationship between predicted value and standardized residuals.

As Figure 1 shows, the residuals of the logprice model are closer to a normal distribution than those of the model with price as the dependent variable, because the line of points in the leftmost two q-q-plots is straighter for logprice. Also when predicting logprice instead of price, predicted values and real values seem to share a more similar distribution, as the middle two q-q-plots show. Also there seems to be a curvilinear relationship between predicted values and residuals when using price directly, which does not seem to be the case when using logprice. Finally also the adjusted R^2 value is higher for the logprice model ($R^2 = 0.830$) than for the price model ($R^2 = 0.802$), which suggests a better prediction. Therefore we settle at the logprice as our dependent variable.

4.3 best subset selection

With our 8 predictor variables, best subset selection is performed to determine which model can predict the logprice best. In total, 255 models ($= \sum_{i=1}^8 \binom{8}{i} = 2^8 - 1$) have been fitted in iterations $i \in \{1, \dots, 8\}$ to the data. Technically the number of predictors k used for models in iteration i was not always the same, since categorical variables were split up into two dichotomous variables each (dummy coding), before performing model fitting. The models of each iteration were then sorted for their BIC and AIC values and the best value from each iteration i can be seen in Table 2

Table 2: Best subset selection: lowest AIC and BIC in each iteration and the respective model predictors added on top of predictors in all table cells above

i	AIC	BIC	predictors	models calculated
1	213.46	225.71	model	8
2	-215.15	-198.82	age	28
3	-414.14	-393.73	mileage	56
4	-564.22	-535.65	transmission	70
5	-651.78	-619.12	engineSize	56
6	-706.25	-665.42	fuelType	28
7	-706.41	-661.51	fuel consumption	8
8	-706.35	-657.37	tax	1

The AIC suggests that the model with 7 predictor variables (all but tax) is the best, while the BIC wants us to stop at iteration 6 and rises again for $i \geq 7$. So the BIC suggests that in addition to tax, also fuel consumption should be excluded from the final model. We will decide for the model the BIC suggests, which will include the predictors: model, age, mileage, transmission, engineSize and fuelType.

5 Summary

Table 3 shows the regression coefficients and their confidence intervals for the final model with 6 predictors (technically 9, if we count the dummy variables as single predictors). The model exhibits an adjusted R^2 value of 0.962 which is quite impressive and signifies a very good model fit. The good model fit can also be seen in the scatter plot in Figure 1 exhibiting a strong linear relationship between predicted logprice and real logprice.




Figure 2: relationship between predicted logprice and real logprice

All regression coefficients differ from 0 significantly since their 95%-confidence intervals $([CI_l, CI_u])$ do not contain 0.

Table 3: The final regression model predicting logprice

	$\hat{\beta}_j$	$SE(\hat{\beta}_j)$	CI_l	CI_u
Intercept	2.6160	0.056	2.506	2.726
age	-0.0884	0.004	-0.096	-0.081
engineSize	0.2865	0.029	0.229	0.344
fuelTypePetrol	0.1227	0.020	0.084	0.161
fuelTypeHybrid	0.4623	0.037	0.389	0.535
mileage	-0.0058	0.000	-0.006	-0.005
modelTRoc	0.1534	0.018	0.118	0.188
modelUp	-0.5180	0.024	-0.566	-0.470
transmissionAutomatic	0.1221	0.019	0.085	0.159
transmissionSemiAuto	0.1180	0.018	0.083	0.153

Bibliography

Exchange & mart: New & used cars for sale near you.
<https://www.exchangeandmart.co.uk/>. (Accessed on 01/16/2022).

Understanding diagnostic plots for linear regression analysis | university of virginia library research data services + sciences.
<https://data.library.virginia.edu/diagnostic-plots/>. (Accessed on 01/16/2022).

Ludwig Fahrmeir, Thomas Kneib, Stefan Lang, and Brian Marx. Regression models. In *Regression*. Springer, 2013.

Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani. *An introduction to statistical learning*, volume 112. Springer, 2013.

Selva Prabhakaran. 10 assumptions of linear regression - full list with examples and code. <http://r-statistics.co/Assumptions-of-Linear-Regression.html>. (Accessed on 01/16/2022).

Appendix

A Additional figures